

# **On-Device AI 실습**

# **Quantization for LLM**

---

Youngmin Jeon

---

# Overview

---

## 1. Weight-only Quantization

1. AWQ

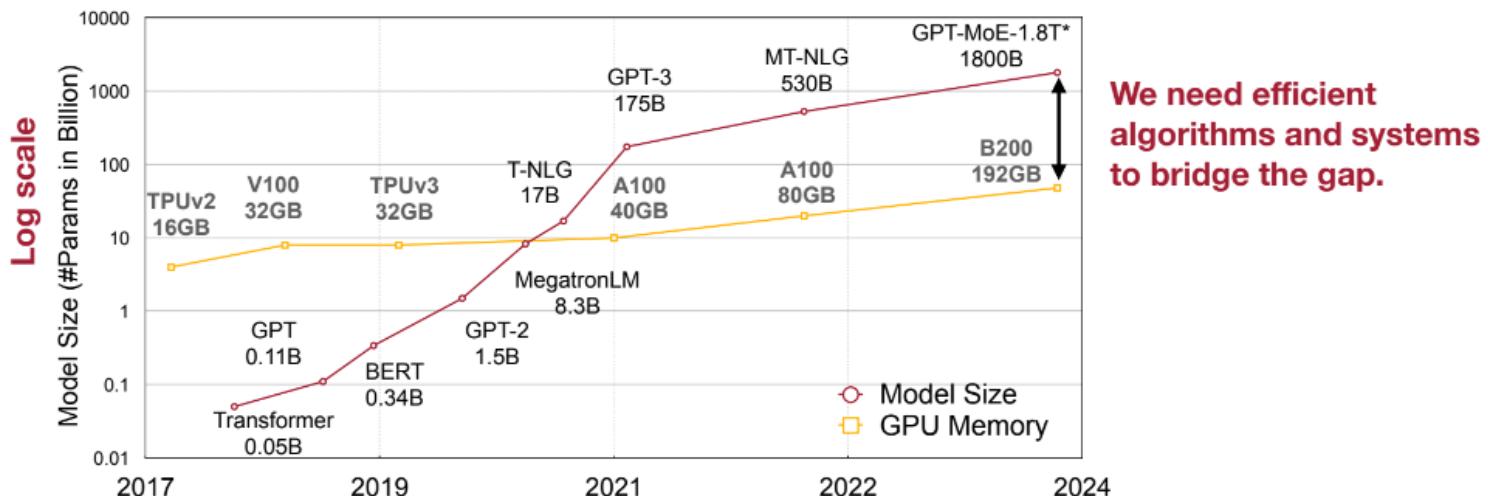
## 2. Weight and Activation Quantization

1. SmoothQuant
2. QuaRot/SpinQuant

# Challenge for LLM deployment

Despite being powerful, LLMs are hard to serve on the edge

- LLM sizes and computation are increasing exponentially.
- Domain-specific accelerator alone is not enough.
  - We need model compression techniques and system support to bridge the gap.



We need efficient  
algorithms and systems  
to bridge the gap.

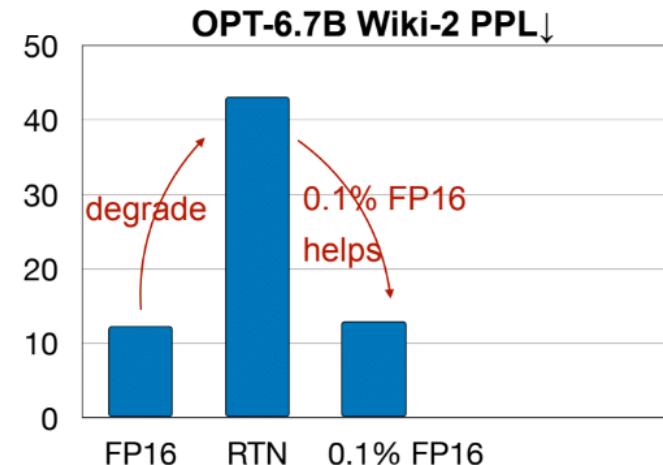
# AWQ: Activation-aware Weight Quantization

**Observation: Weights are not equally important; 0.1% salient weights**

$W_{FP16}$				$Q(W)_{MixPrec}$			
+1.2	-0.2	-2.4	-3.4	+1	+0	-2	-3
-2.5	-3.5	+1.9	+1.4	-2.5	-3.5	+1.9	+1.4
-0.9	+1.6	-2.5	-1.9	-1	+2	-3	-2
-3.5	+1.5	+0.5	-0.1	-4	+2	+1	+0
+1.8	-1.6	-3.2	-3.4	+2	-2	-3	-3
+2.4	-3.5	-2.8	-3.9	+2	-4	-3	-4
+0.1	-3.8	+2.4	+3.4	+0	-4	+2	+3
+0.9	+3.3	-1.9	-2.3	+1	+3	-2	-2

RTN

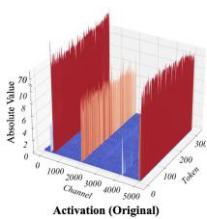
FP16  
channel



- We find that weights are not equally important, keeping **only 0.1%** of salient weight channels in FP16 can greatly improve perplexity
- But how do we select salient channels? Should we select based on weight magnitude?

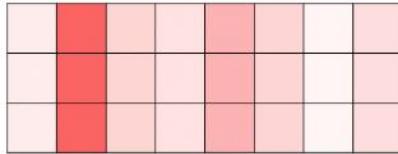
# Protect 1% salient channels

Salient weights are determined by activation distribution, not weight



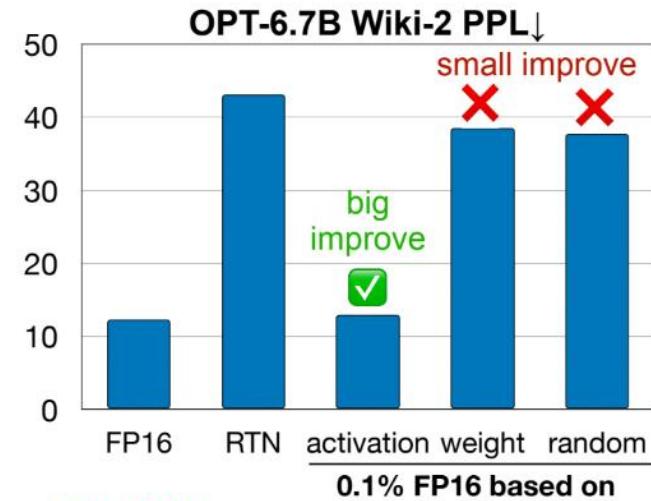
determine the salient weights by activation

$X$



$Q(W)^{\text{MixPrec}}$			
+1	+0	-2	-3
-2.5	-3.5	+1.9	+1.4
-1	+2	-3	-2
-4	+2	+1	+0
+2	-2	-3	-3
+2	-4	-3	-4
+0	-4	+2	+3
+1	+3	-2	-2

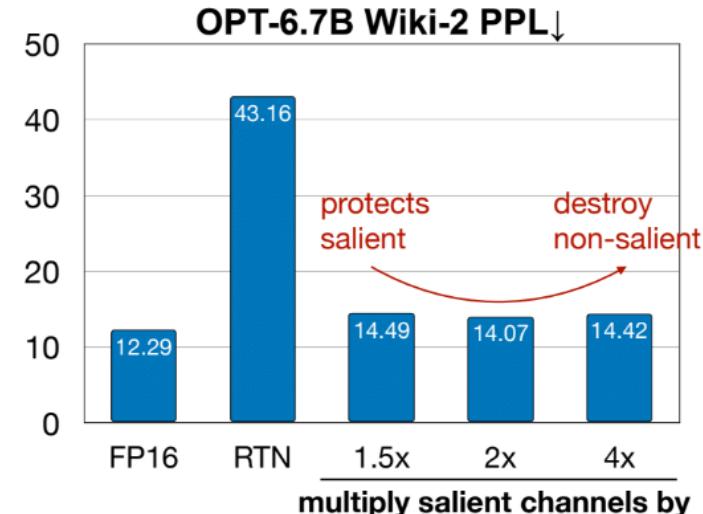
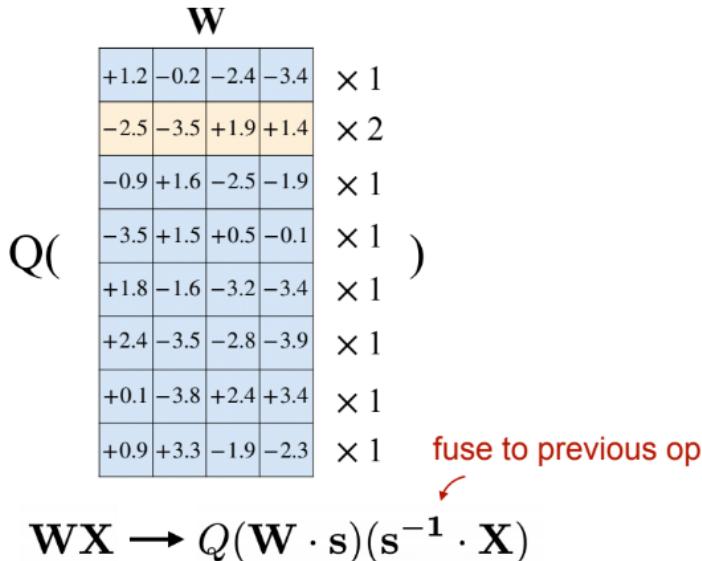
FP16 channel



- We find that weights are not equally important, keeping **only 0.1%** of salient weight channels in FP16 can greatly improve perplexity
- But how do we select salient channels? Should we select based on weight magnitude?
- No! We should look for **activation distribution, but not weight!**

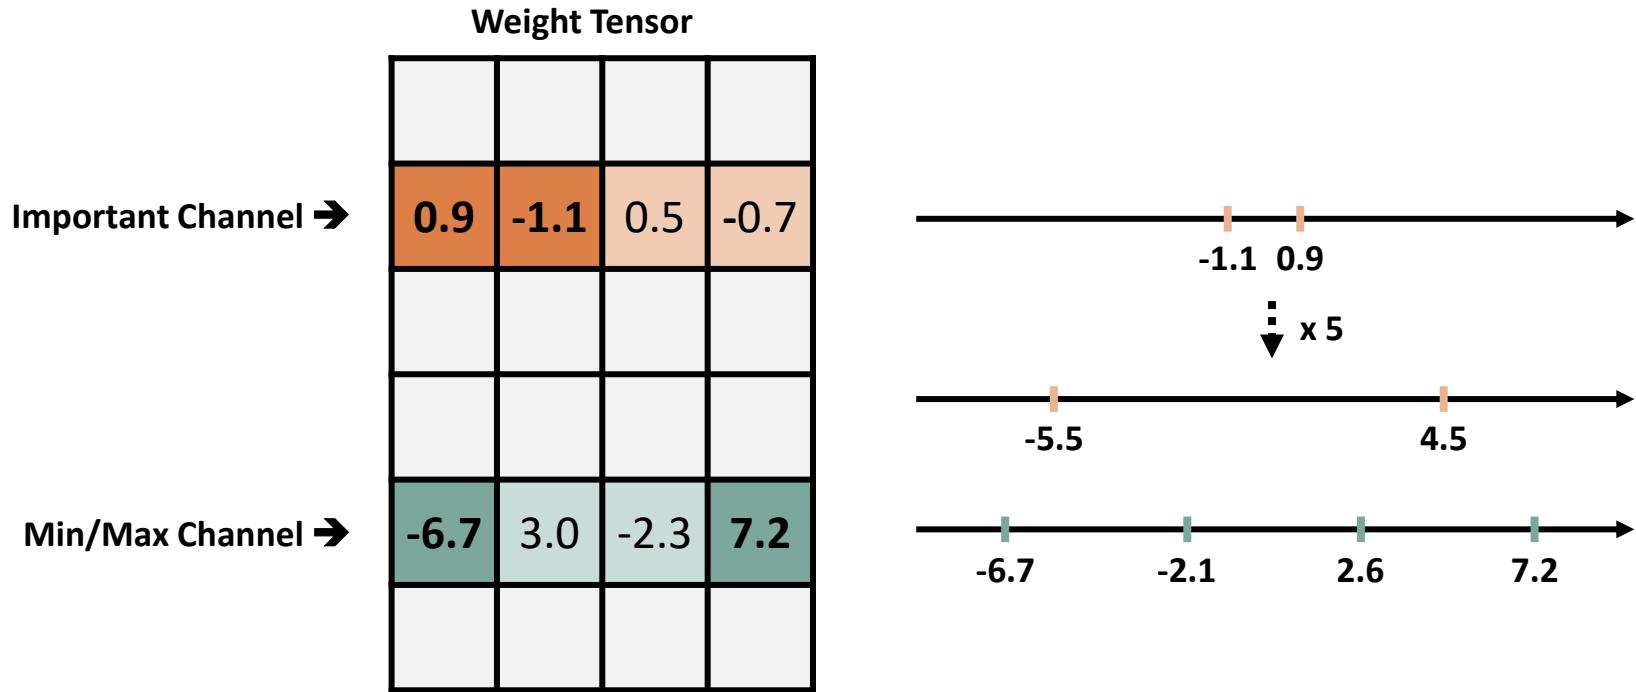
# Scale 0.1% salient channels

## Protecting salient weights by scaling (no mixed prec.)



- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Why?

# Scale salient channels



# Scale 1% salient channels

- **Protecting salient weights by scaling**

- $y = \mathbf{W}\mathbf{x} \rightarrow \tilde{\mathbf{W}} = \Delta \cdot \text{Round}\left(\frac{\mathbf{W}}{\Delta}\right), \Delta = \frac{\max(\mathbf{W})}{2^{N-1}}$

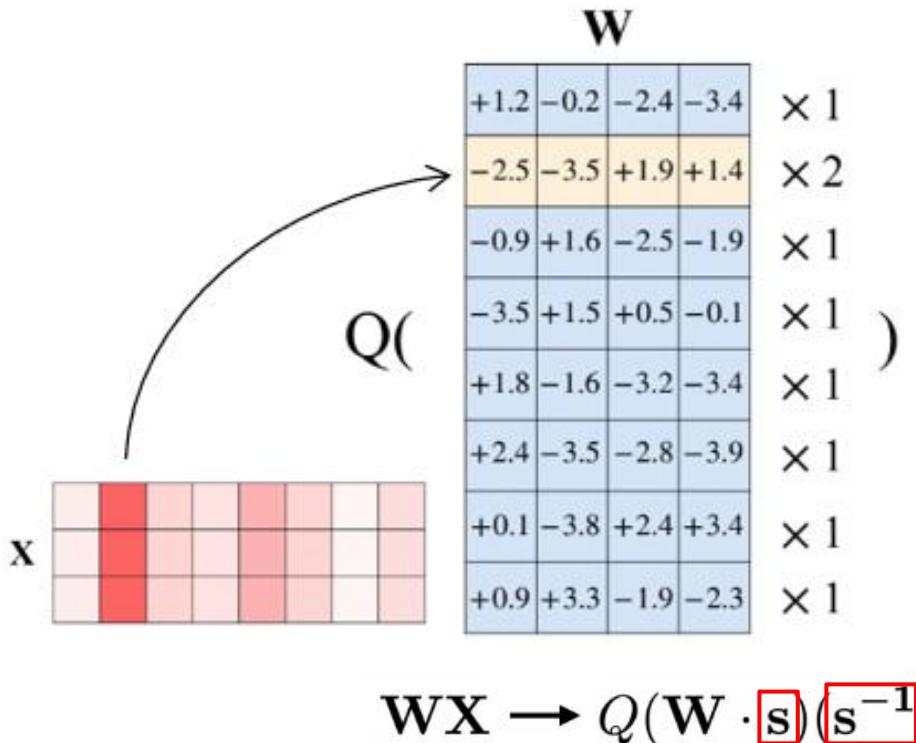
- Error =  $\Delta \cdot \text{RoundError}$

- $y = (\mathbf{W} \cdot s)(\mathbf{x}/s) \rightarrow \widetilde{\mathbf{W} \cdot s} = \Delta' \cdot \text{Round}\left(\frac{\mathbf{W} \cdot s}{\Delta'}\right)$

- $\text{Error}' = \Delta' \cdot \text{RoundError} \cdot \frac{1}{s} = \Delta \cdot \text{RoundError} \cdot \frac{1}{s} = \text{Error} \cdot \frac{1}{s}$

- $\max(\mathbf{W}) = \max(\mathbf{W} \cdot s) \rightarrow \Delta' = \Delta$

# [실습1] Scale 1% salient channels



# Answer

```
##### YOUR CODE STARTS HERE #####
# Step 1: importance를 기준으로 1%의 중요한 채널을 찾으세요 (hint: use torch.topk())
# hint : torch.topk() 함수를 사용하세요. torch.topk() 함수는 PyTorch에서 텐서의 값 중 상위 k개의 값과 그들의 인덱스를 반환하는 함수입니다.
outlier_mask = torch.topk(importance, int(len(importance) * 0.01))[1]
assert outlier_mask.dim() == 1

##### YOUR CODE ENDS HERE #####
# 스케일 팩터를 적용하는 것을 시뮬레이션하기 위해, 양자화 전에 스케일 팩터를 곱하고, 양자화 후에 스케일 팩터로 나눕니다.
# scale_factor를 이용해 중요한 가중치 채널의 값을 확대합니다.
m.weight.data[:, outlier_mask] *= scale_factor

m.weight.data = pseudo_quantize_tensor(m.weight.data, n_bit=m_bit, q_group_size=q_group_size)

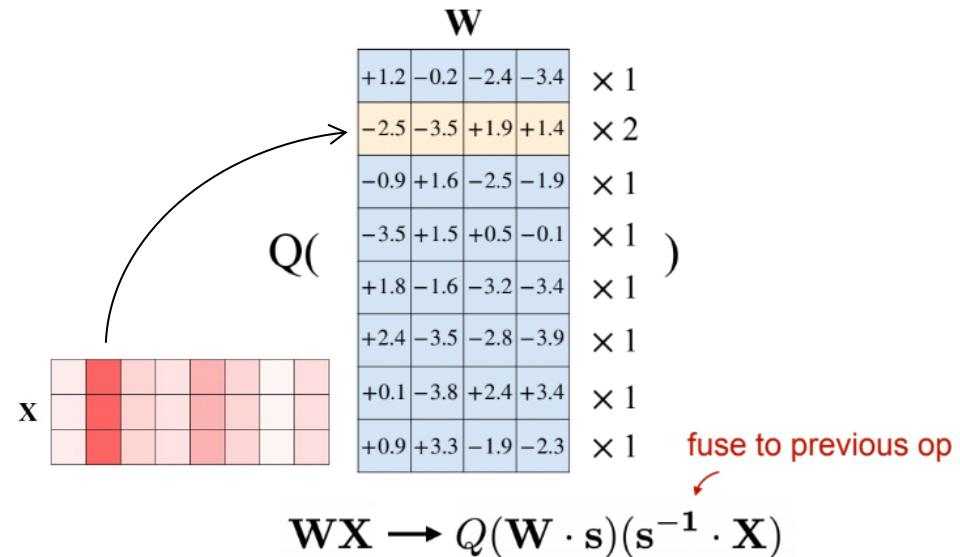
##### YOUR CODE STARTS HERE #####
# Step 2: scale_factor를 이용해 중요한 가중치 채널의 값을 다시 축소하세요.
m.weight.data[:, outlier_mask] /= scale_factor

##### YOUR CODE ENDS HERE #####

```

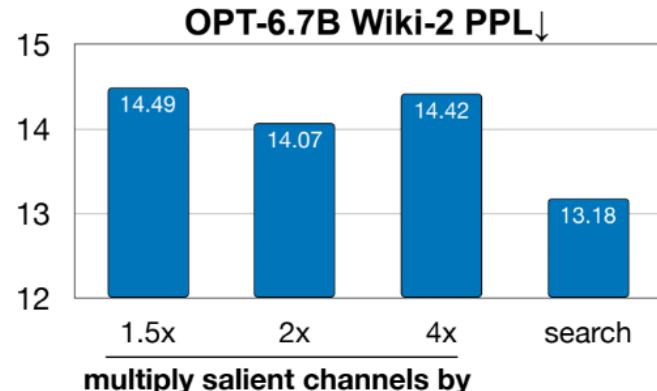
# Scale Factor Search

Protecting salient weights by scaling (no mixed prec.)



- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Take a data-driven approach with a fast **grid search**

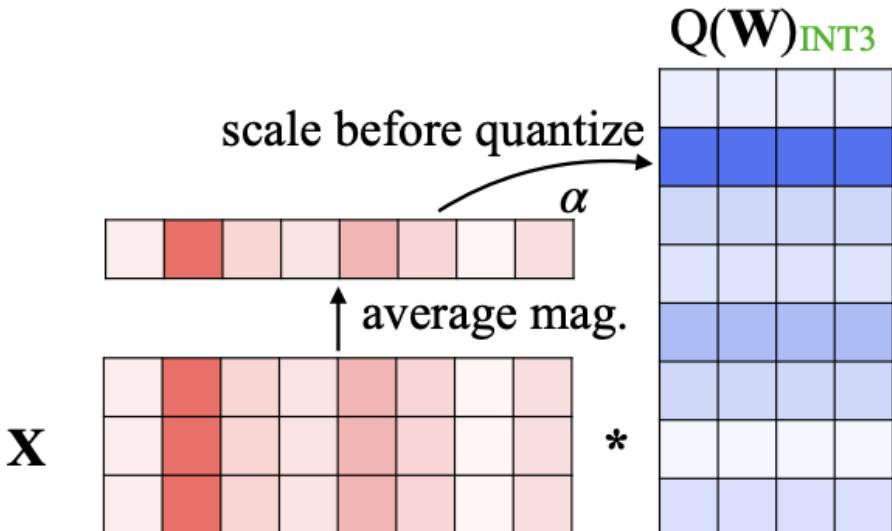
**L1 norm of Activation**



$$\mathcal{L}(\mathbf{s}) = \|Q(\mathbf{W} \cdot \mathbf{s})(\mathbf{s}^{-1} \cdot \mathbf{X}) - \mathbf{WX}\|$$

$\mathbf{s} = \mathbf{s}_X^\alpha$  Activation-awareness is important,  
but not weight-awareness

# [실습2] Scale Factor Search



$$\mathbf{s}_{\mathbf{x}^\alpha} = \|\mathbf{X}\|_1, \quad \|\mathbf{X}\|_1 = \sum |X_i|$$

$\mathbf{s} = \mathbf{s}_{\mathbf{x}^\alpha}$  Activation-awareness is important,  
but not weight-awareness

- Multiplying the salient channels with  $s > 1$  reduces its quantization error
- Take a data-driven approach with a fast **grid search**

# Answer

```
##### YOUR CODE STARTS HERE #####
# Step 2: 공식에 따라 스케일 계산: scales = s_x^ratio
scales = s_x ** ratio # must clip the s_x, otherwise will get nan later

assert scales.shape == s_x.shape

##### YOUR CODE ENDS HERE #####
scales = scales / (scales.max() * scales.min()).sqrt().view(1, -1)

for fc in linears2scale:

    scales = scales.to(fc.weight.device)

    # scale_factor를 이용해 중요한 가중치 채널의 값을 확대합니다.
    fc.weight.mul_(scales)

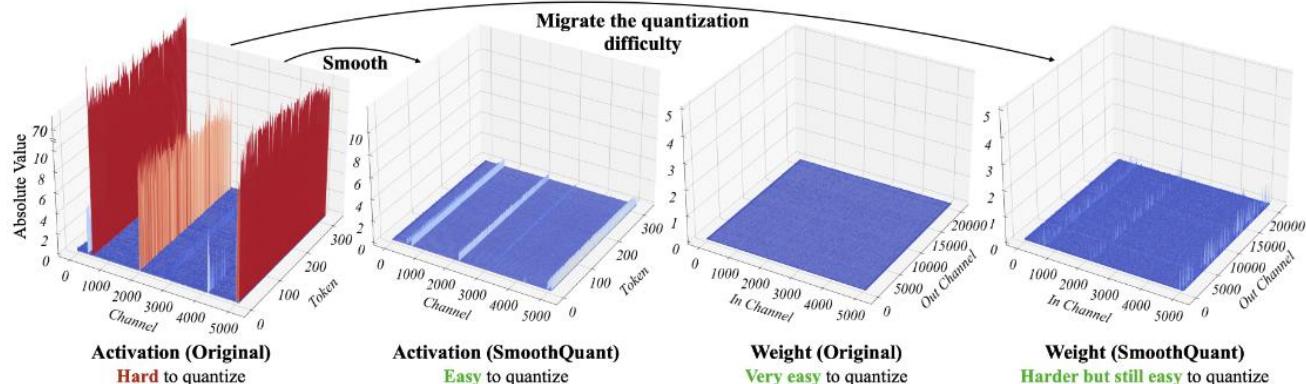
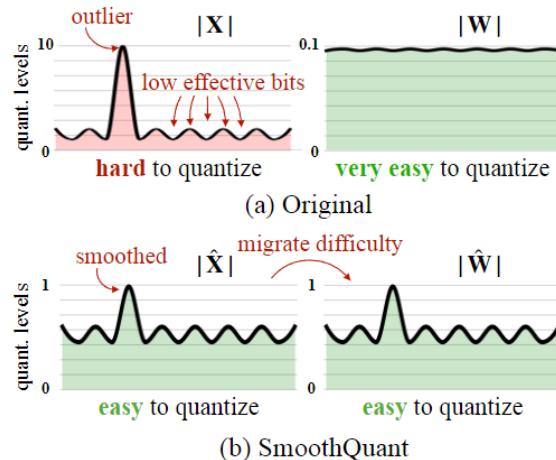
    fc.weight.data = pseudo_quantize_tensor(fc.weight.data, w_bit, q_group_size)

##### YOUR CODE STARTS HERE #####
# Step 3: scale_factor를 이용해 중요한 가중치 채널의 값을 다시 축소하세요.
fc.weight.data /= scales

##### YOUR CODE ENDS HERE #####
```

# SmoothQuant

- SmoothQuant's intuition
  - Migrate scale from activations to weights  $W$  before quantization



# SmoothQuant

- Main idea of SmoothQuant

$$\mathbf{Y} = \mathbf{X} \cdot \mathbf{W}, \mathbf{Y} \in \mathbb{R}^{T \times C_o}, \mathbf{X} \in \mathbb{R}^{T \times C_i}, \mathbf{W} \in \mathbb{R}^{C_i \times C_o},$$

$$s_j = \max(|\mathbf{X}_j|)^\alpha / \max(|\mathbf{W}_j|)^{1-\alpha}$$

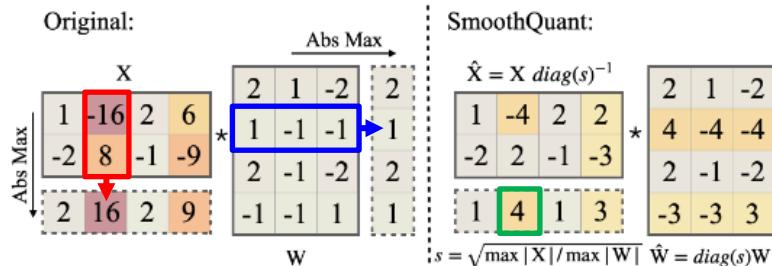


Figure 5: Main idea of SmoothQuant when  $\alpha$  is 0.5. The smoothing factor  $s$  is obtained on calibration samples and the entire transformation is performed offline. At runtime, the activations are smooth without scaling.

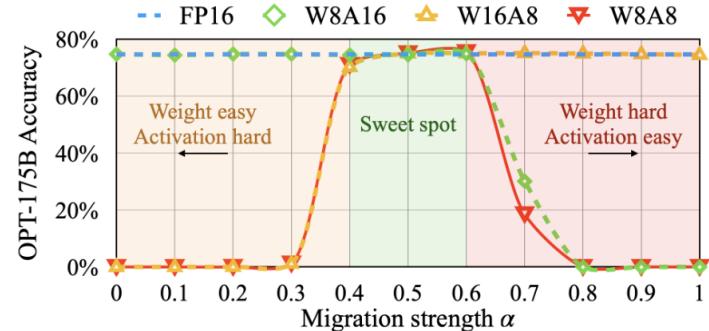
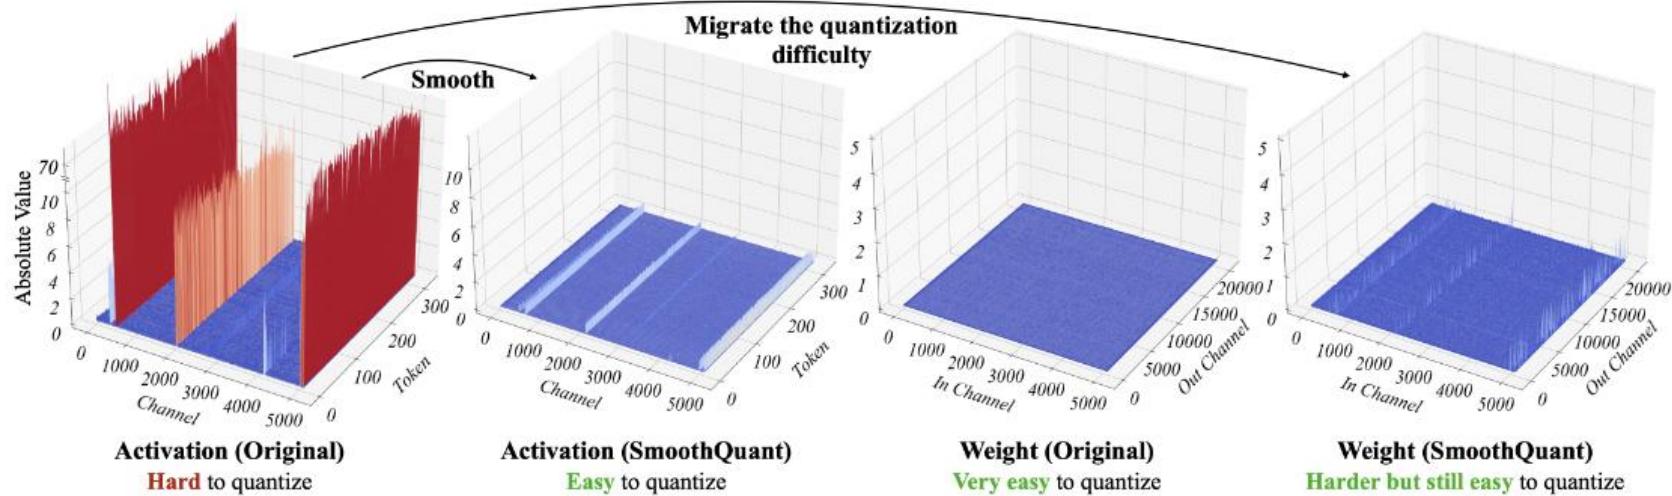


Figure 10: A suitable migration strength  $\alpha$  (sweet spot) makes both activations and weights easy to quantize. If the  $\alpha$  is too large, weights will be hard to quantize; if too small, activations will be hard to quantize.

# [실습3] Quantization Difficulty Migration



# Answer

17

23

```
def smooth_ln_fcs_by_scale(ln, fcs, scale):
    if not isinstance(fcs, list):
        fcs = [fcs]
    assert isinstance(ln, nn.LayerNorm)
    for fc in fcs:
        assert isinstance(fc, nn.Linear)
    ##### YOUR CODE STARTS HERE #####
    # Step 1: layernorm의 weight와 bias를 scale로 나누어주세요. (hint: div_()함수를 통해 tensor 전체를 특정한 값으로 나누어 줄 수 있습니다.)
    ln.weight.div_(scale)
    ln.bias.div_(scale)
    ##### YOUR CODE ENDS HERE #####
    for fc in fcs:
        ##### YOUR CODE STARTS HERE #####
        # Step 2: fc의 weight에 scale을 곱해주세요. (hint: mul_()함수를 통해 tensor 전체에 특정한 값을 곱해 줄 수 있습니다.)
        fc.weight.mul_(scale)
        ##### YOUR CODE ENDS HERE #####
```

# [실습4] Scale Factor Search

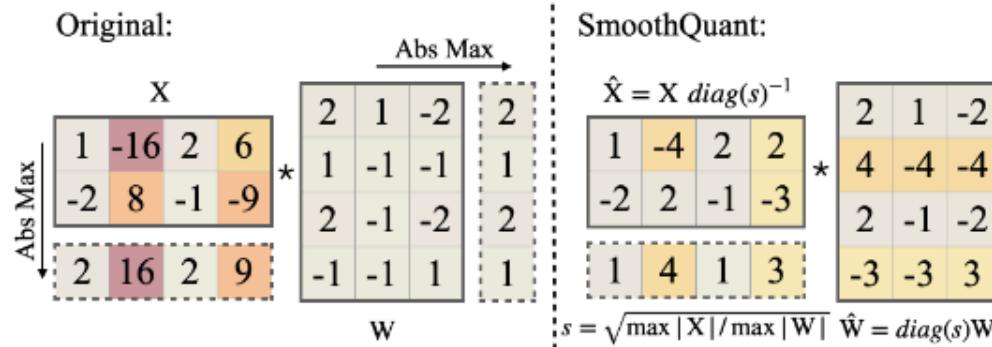


Figure 5: Main idea of SmoothQuant when  $\alpha$  is 0.5. The smoothing factor  $s$  is obtained on calibration samples and the entire transformation is performed offline. At runtime, the activations are smooth without scaling.

# Answer

```
scales = (
    ##### YOUR CODE STARTS HERE #####
    #Activation Scales 값과 Weight Scales 값에 alpha를 적절히 거듭제곱해주어야 합니다.
    #Hint: pow()함수를 통해서 거듭제곱을 사용할 수 있습니다.
    (act_scales.pow(alpha) / weight_scales.pow(1 - alpha))
    ##### YOUR CODE ENDS HERE #####
)
```

# Ratation based Quantization

$$\mathbf{Y} = (\mathbf{X}\mathbf{R})(\mathbf{R}^{-1}\mathbf{W}^{-1}) = \mathbf{X}\mathbf{W}^T$$

A rotation matrix is an orthogonal matrix  $\mathbf{R}$  satisfied  $\mathbf{R}\mathbf{R}^T = 1$  and  $|\mathbf{R}| = 1$

## QuaRot

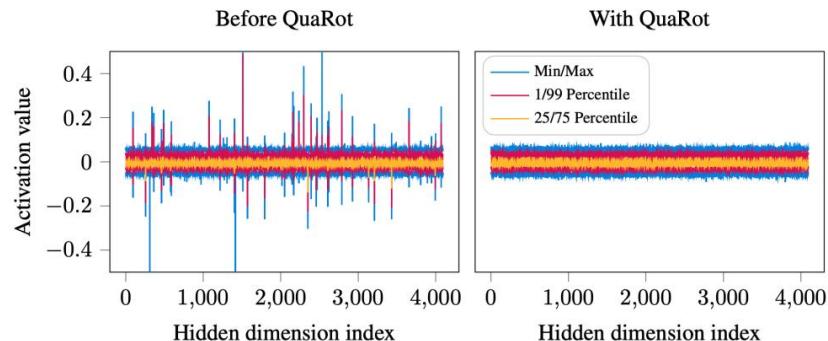
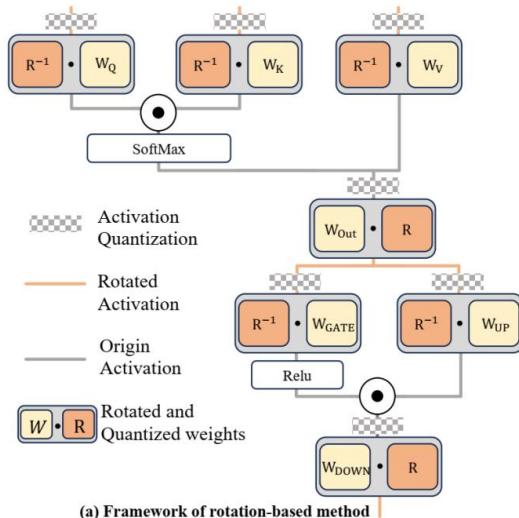
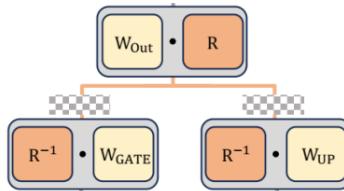


Figure 1: The distributions of activations at the input to the FFN block in LLaMA2-7B model, in the tenth layer. Left: using the default configuration as downloaded from Hugging Face. Right: after processing using QuaRot. The processed distribution has no outliers, leading to superior quantization.

# Layernorm $\Leftrightarrow$ Linear Fusion

- $y = x \cdot W_{out} \cdot R \cdot R^{-1} \cdot W_{GATE} = x \cdot W_{out} \cdot W_{GATE}$

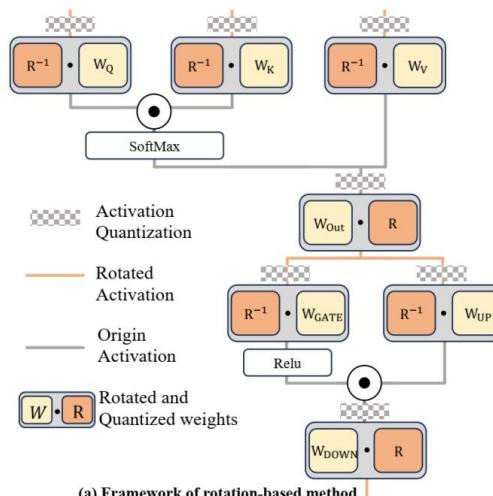


- In real case,  $y = LN(x \cdot W_{out} \cdot R) \cdot R^{-1} \cdot W_{GATE} \neq x \cdot W_{out} \cdot W_{GATE}$ 
  - $LN(x) = x \cdot \gamma + \beta$
  - $LN(x) \cdot W + b = (x \cdot \gamma + \beta) \cdot W + b$   
 $= x \cdot \gamma W + \beta W + b = x \cdot \tilde{W} + \tilde{b}$
  - $y = LN(x \cdot W_{out} \cdot R) \cdot R_{-1} \cdot \widetilde{W_{GATE}} = x \cdot W_{out} \cdot \widetilde{W_{GATE}}$

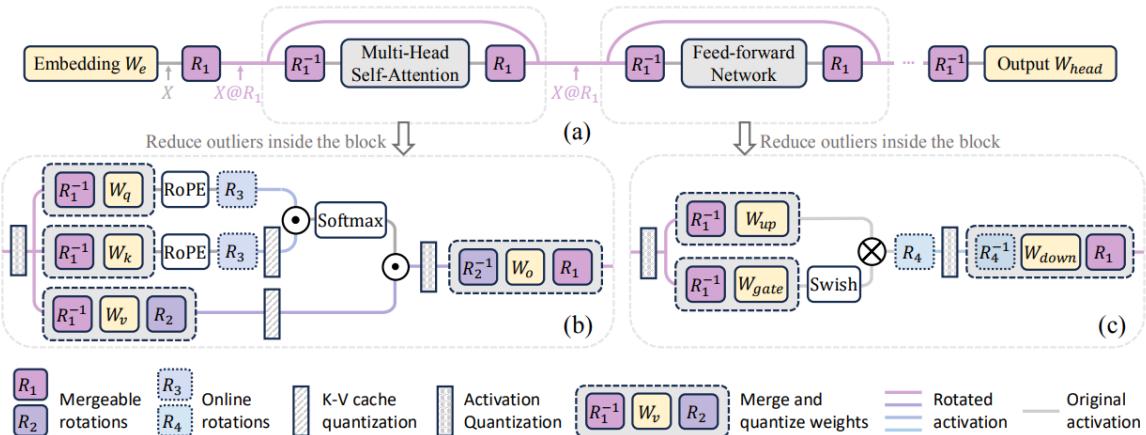
# [실습5] Rotate Matrix 적용

- nn.Linear 연산은  $W^T$  형태로 저장
  - $W^T \cdot R \rightarrow$  연산시에  $R^T \cdot W$
- Embedding Parameter Shape :  
(Num\_Tokens, Hidden\_dim)
- Linear Parameter Shape :  
(Output\_Channel, Input\_Channel)
- Roation Matrix Shape :  
(Hidden\_dim, Hidden\_dim)

## QuaRot



## SpinQuant



# Answer

```
##### YOUR CODE STARTS HERE #####
# Pytorch에서 @ 연산이 Dot Product 임을 사용하시기 바랍니다.
# nn.Linear 연산의 Parameter는 W^T 형태로 저장되어 있다는 것을 유의하시기 바랍니다.
# Embedding Parameter Shape : (Num_Tokens, Hidden_dim)
# Linear Parameter Shape : (Output_Channel, Input_Channel)
# Roation Matrix Shape : (Hidden_dim, Hidden_dim)

if isinstance(m, nn.Embedding):
    W_ = m.weight.data
    m.weight.data = W_ @ R1

if isinstance(m, nn.Linear):
    if "out_proj" in n or "fc2" in n:
        # Att Out Proj, FFN Down Proj
        W_ = m.weight.data
        m.weight.data = R1.T @ W_
    else:
        # QKV Proj, FFN Up Proj, FFN Gate Proj
        W_ = m.weight.data
        m.weight.data = W_ @ R1
    ➔ x · W · R

else:
    # QKV Proj, FFN Up Proj, FFN Gate Proj
    W_ = m.weight.data
    m.weight.data = W_ @ R1
    ➔ x · RT · W

##### YOUR CODE ENDS HERE #####
```